

Images

Jenny Hong Ahmed Bou-Rabee Stephen Boyd

EE103
Stanford University

October 18, 2015

Outline

Representation

Linear operations

In-painting

Image de-blurring

Monochrome images

- ▶ a.k.a. *monochrome* or *gray-scale* image
- ▶ image represented by its brightness values at an array of $m \times n$ locations (pixels)
- ▶ typical sizes
 - thumbnail: 16×16 , 64×64 , 128×128
 - 4K \times 6K = 24M pixels
 - HD is 1280×720 pixels
- ▶ can represent by an $m \times n$ matrix X or by a single vector $x \in \mathbf{R}^{mn}$ with some encoding of the pixel locations, e.g.,

$$X_{ij} = x_k, \quad k = m(j-1) + i, \quad k = 1, \dots, mn$$

(this stacks the columns of X , from left to right)

Brightness values

- ▶ x_i is brightness of pixel i
- ▶ typically $0 \leq x_i \leq 1$ where 0 is black and 1 is white
- ▶ values outside $[0, 1]$ are clipped (so $x_i < 0$ shows up as black)
- ▶ if x is an image, $-x$ is completely black
- ▶ *negative image* is given by $1 - x$
- ▶ $\text{avg}(x)$ is average intensity (brightness) of image
- ▶ $\text{std}(x)$ corresponds to image *contrast*

Scaling, shifting, and adding images

- ▶ what does image

$$y = a(x - \mathbf{avg}(x)\mathbf{1}) + (\mathbf{avg}(x) + b)\mathbf{1} = ax + c\mathbf{1}$$

$(c = (1 - a) \mathbf{avg}(x) + b)$ look like?

- a scale contrast
- b shifts brightness

- ▶ $y_i = x_i^\gamma$ is called γ -correction (widely used)
- ▶ if x and y are images, $x + y$ is perceived as composite or combination of the images (and isn't natural, except in some cases)

Examples

original



original + 0.5



$(\text{original} - 0.4) * 10$



Examples

pumpkins



flowers



$(\text{pumpkins} + \text{flowers}) / 2$



Color images

- ▶ humans perceive 3 colors, which can be represented in different ways (e.g., RGB, CMYK)
- ▶ most common is RGB (Red-Green-Blue)
- ▶ color represented as a 3-vector (r, g, b) , with r, g, b between 0 and 1
 - $(1, 0, 0)$ is bright red
 - $(1, 0, 1)$ is bright purple
 - $(0.2, 0.2, 0.2)$ is a gray
- ▶ $m \times n$ image given by 3 $m \times n$ matrices or one vector $x \in \mathbf{R}^{3mn}$

Colors


- ▶ $(1,0,0)$, $(0,1,0)$, $(0,0,1)$
- ▶ $(1,1,0)$, $(0,1,1)$, $(1,0,1)$
- ▶ $(0.2,0.2,0.2)$, $(0.5,0.5,0.5)$, $(0.75,0.75,0.75)$

Color images

original




red



green



blue



Converting color to monochrome


- ▶ color pixel values converted to monochrome using $y_i = w^T(r_i, g_i, b_i)$
 - obvious choice: $w = (1/3, 1/3, 1/3)$
 - another common choice: $w = (0.299, 0.587, 0.114)$
 - other choices used for special effects

Converting color to monochrome

original




equal weights



weights: 0.2, 0.5, 0.3



weights: 0.6, -0.4, 0.8



Video

- ▶ video is represented as a sequence of images captured periodically
- ▶ each image is called a *frame*
- ▶ typical frame rates: 24, 30, or 60 frames per second

Outline

Representation

Linear operations

In-painting

Image de-blurring

Linear image mappings

- ▶ for y and x images, linear mapping $y = Ax$ can represent many common operations on images
 - color to monochrome conversion
 - color correction
 - any mapping from original to distorted pixel locations (e.g., flipping, stretching)
 - blurring
 - changing to lower or higher resolution
 - vertical and horizontal differencing

Moving pixels

- ▶ pixel at location i in image y is the pixel at value $j = d(i)$ in image x
- ▶ $d(i)$ gives *distortion map*
- ▶ examples: flipping, zooming, rotating, shifting, key correction
- ▶ some issues/details:
 - we'll need to approximate the location of the pixels
 - we need to do something with y pixels that don't correspond to any x pixels
- ▶ $y = Ax$, where i th row of A is $e_{d(i)}^T$
(or 0, if y_i doesn't correspond to any x pixel)

Flipping images

original image



horizontal flip



vertical flip



Blurring images

- ▶ represent image as $m \times n$ matrix X
- ▶ represent blur *point spread function* as $p \times q$ matrix B
- ▶ blurred image is given by Y with

$$Y_{ij} = \sum_{k,l} X_{i-k+1,j-l+1} B_{k,l}$$

where

- the sum is over all integers k, l
- we interpret X_{ij} and $B_{k,l}$ as zero when the indices are out of range
- ▶ called *2-D convolution* of X and B , denoted $Y = X * B$ or $Y = A \star B$
- ▶ blurring is model of effects of optical imperfections, motion blur, ...

Blurring images


original image



blurred image



point spread function



Horizontal and vertical differences

- ▶ X is $m \times n$ image (matrix), x its mn -vector representation
- ▶ horizontal first order difference is $m \times (n - 1)$ matrix Y with

$$Y_{ij} = X_{i,j+1} - X_{i,j}, \quad i = 1, \dots, m, \quad j = 1, \dots, n - 1$$

- ▶ vertical first order difference is $(m - 1) \times n$ matrix Z with

$$Z_{ij} = X_{i+1,j} - X_{i,j}, \quad i = 1, \dots, m - 1, \quad j = 1, \dots, n$$

- ▶ these are linear operations, so we have

$$y = D^{\text{horiz}}x, \quad z = D^{\text{vert}}x$$

for an $m(n - 1)$ -matrix D^{horiz} and an $(m - 1)n$ -matrix D^{vert}

- ▶ each row contains one +1 and one -1

Horizontal and vertical differences

(shown for 3×3 image)


$$D^{\text{horiz}} = \begin{bmatrix} -1 & & & +1 & & & & \\ & -1 & & & +1 & & & \\ & & -1 & & & +1 & & \\ & & & -1 & & & +1 & \\ & & & & -1 & & & +1 \\ & & & & & -1 & & \\ & & & & & & -1 & \\ & & & & & & & +1 \end{bmatrix}$$
$$D^{\text{vert}} = \begin{bmatrix} -1 & & & & & & & \\ & +1 & & & & & & \\ & & -1 & & & & & \\ & & & +1 & & & & \\ & & & & -1 & & & \\ & & & & & +1 & & \\ & & & & & & -1 & \\ & & & & & & & +1 \end{bmatrix}$$

Horizontal and vertical differences


original image



horizontal difference



vertical difference



Laplacian

- ▶ the *Laplacian* function is

$$\begin{aligned}\mathcal{L}(x) &= \|D^{\text{horiz}}x\|^2 + \|D^{\text{vert}}x\|^2 \\ &= \sum_{i=1}^{m-1} \sum_{j=1}^{n-1} ((X_{i+1,j} - X_{i,j})^2 + (X_{i,j+1} - X_{i,j})^2)\end{aligned}$$

(we also write $\mathcal{L}(X)$)

- ▶ $\mathcal{L}(X)$ is a measure of roughness of the image X
 - $\mathcal{L}(X)$ is small when the image is smooth
 - $\mathcal{L}(X) = 0$ only if the image is constant
- ▶ $\mathcal{L}(X)$ is used as a *regularizer*

Outline

Representation

Linear operations

In-painting

Image de-blurring

In-painting

- ▶ we are given an image with some pixels values unknown
- ▶ *in-painting* means to guess values of the unknown pixels so the recovered image looks good or natural
- ▶ in example below, unknown values are shown as black



Least-squares in-painting


- ▶ corrupted/damaged image is given by $m \times n$ matrix X^{corr}
- ▶ $\mathcal{K} \subset \{1, \dots, m\} \times \{1, \dots, n\}$ are the indices of known pixels
- ▶ we need to choose an image X that agrees with the given image on known pixels: $X_{ij} = X_{ij}^{\text{corr}}, (i, j) \in \mathcal{K}$
- ▶ we'll choose X to minimize $\mathcal{L}(X)$, the sum square deviation of all pixel values from their neighbors (small $\mathcal{L}(X)$ gives a smooth image)
- ▶ a least-squares problem (variables are unknown pixel values)

In-painting


original image



damaged image



inpainted image



Outline

Representation

Linear operations

In-painting

Image de-blurring

Corrupted image

- ▶ y is a linear function of x^{true} , with noise:

$$y = Ax^{\text{true}} + v$$

- ▶ y is the *corrupted* image, which we have
- ▶ x^{true} is the original image, which we want to guess/recover
- ▶ v is a noise, which we assume is small
- ▶ A is a (known) matrix, often a blurring operator
- ▶ *image de-blurring* is guessing x^{true}
- ▶ even if A is invertible, the guess $x = A^{-1}y$ could look very bad

Least-squares de-blurring

- ▶ *least-squares de-blurring*: choose x to minimize

$$\|Ax - y\|^2 + \lambda \mathcal{L}(x)$$

- ▶ first term is $\|v\|^2$
- ▶ $\lambda > 0$ is a regularization parameter
 - large λ makes x smooth
 - small λ makes $\|Ax - y\|^2$ small

De-blurring example

original image



corrupted image




deblurred image with $\lambda = 0.03$



De-blurring: Effect of regulariziton

$\lambda = 0.0003$



$\lambda = 0.03$



$\lambda = 3$

